

Assignment 10: Data Scraping

Xia Meng Howey

OVERVIEW

This exercise accompanies the lessons in Environmental Data Analytics on data scraping.

Directions

1. Rename this file <FirstLast>_A10_DataScraping.Rmd (replacing <FirstLast> with your first and last name).
2. Change “Student Name” on line 3 (above) with your name.
3. Work through the steps, **creating code and output** that fulfill each instruction.
4. Be sure your code is tidy; use line breaks to ensure your code fits in the knitted output.
5. Be sure to **answer the questions** in this assignment document.
6. When you have completed the assignment, **Knit** the text and code into a single PDF file.

Set up

1. Set up your session:
 - Load the packages `tidyverse`, `rvest`, and any others you end up using.
 - Check your working directory

```
#1  
library(tidyverse)  
library(rvest)  
  
getwd()
```

```
## [1] "/home/guest/R/EDE_Fall2025"
```

2. We will be scraping data from the NC DEQs Local Water Supply Planning website, specifically the Durham’s 2024 Municipal Local Water Supply Plan (LWSP):
 - Navigate to <https://www.ncwater.org/WUDC/app/LWSP/search.php>
 - Scroll down and select the LWSP link next to Durham Municipality.
 - Note the web address: <https://www.ncwater.org/WUDC/app/LWSP/report.php?pwsid=03-32-010&year=2024>

Indicate this website as the URL to be scraped. (In other words, read the contents into an `rvest` webpage object.)

```
#2
website <- read_html(
  'https://www.ncwater.org/WUDC/app/LWSP/report.php?pwsid=03-32-010&year=2024')
website

## {html_document}
## <html xmlns="http://www.w3.org/1999/xhtml" lang="en" xml:lang="en">
## [1] <head>\n<title>DWR :: Local Water Supply Planning</title>\n<meta http-equiv ...
## [2] <body id="plan">\r\n<!--<div id="division-header">\r\n<a name="top" href= ...
```

3. The data we want to collect are listed below:

- From the “1. System Information” section:
 - Water system name
 - PWSID
 - Ownership
- From the “3. Water Supply Sources” section:
 - Maximum Day Use (MGD) - for each month

In the code chunk below scrape these values, assigning them to four separate variables.

HINT: The first value should be “Durham”, the second “03-32-010”, the third “Municipality”, and the last should be a vector of 12 numeric values (represented as strings)“.

```
#3
systemname <- website %>%
  html_nodes("div+ table tr:nth-child(1) td:nth-child(2)") %>%
  html_text()
systemname
```

```
## [1] "Durham"
```

```
PWSID <- website %>%
  html_nodes("td tr:nth-child(1) td:nth-child(5)") %>%
  html_text()
PWSID
```

```
## [1] "03-32-010"
```

```
ownership <- website %>%
  html_nodes("div+ table tr:nth-child(2) td:nth-child(4)") %>%
  html_text()
ownership
```

```
## [1] "Municipality"
```

```

MGD <- website %>%
  html_nodes("th~ td+ td") %>%
  html_text()
MGD

## [1] "34.5000" "36.0600" "37.3300" "32.1000" "46.6500" "37.3600" "38.2000"
## [8] "41.9000" "36.5800" "36.7300" "42.9600" "34.4500"

```

- Convert your scraped data into a dataframe. This dataframe should have a column for each of the 4 variables scraped and a row for the month corresponding to the withdrawal data. Also add a Date column that includes your month and year in data format. (Feel free to add a Year column too, if you wish.)

TIP: Use `rep()` to repeat a value when creating a dataframe.

NOTE: It's likely you won't be able to scrape the monthly withdrawal data in chronological order. You can overcome this by creating a month column manually assigning values in the order the data are scraped: "Jan", "May", "Sept", "Feb", etc...

- Create a line plot of the maximum daily withdrawals across the months for 2024, making sure, the months are presented in proper sequence.

```


#4
df.Durham.water <- data.frame("Month" = c("Jan", "May", "Sep",
                                             "Feb", "Jun", "Oct",
                                             "Mar", "Jul", "Nov",
                                             "Apr", "Aug", "Dec"),
                               "Year" = rep(2024,12),
                               "System Name" = systemname,
                               "PWSID" = PWSID,
                               "Ownership" = ownership,
                               "MGD" = as.numeric(MGD))

df.Durham.water <- df.Durham.water %>%
  mutate(Date = my(paste(Month, "-", Year))) %>%
  arrange(Date)

#5
ggplot(df.Durham.water, aes(
  x=Date,
  y=MGD)) +
  geom_line() +
  labs(title = paste("2024 Max Daily Withdrawals in Durham"))

```

2024 Max Daily Withdrawals in Durham



6. Note that the PWSID and the year appear in the web address for the page we scraped. Construct a function with two inputs - “PWSID” and “year” - that:

- Creates a URL pointing to the LWSP for that PWSID for the given year
- Creates a website object and scrapes the data from that object (just as you did above)
- Constructs a dataframe from the scraped data, mostly as you did above, but includes the PWSID and year provided as function inputs in the dataframe.
- Returns the dataframe as the function’s output

```
#6.
#Construct the scraping web address
PWSID2 <- '03-32-010'
year2 <- 2024

#create scraping function
scrape.it <- function(PWSID2, year2){

  #retrieve website contents
  website.auto <- read_html(paste0(
    'https://www.ncwater.org/WUDC/app/LWSP/report.php', '?pwsid=', PWSID2,
    '&year=', year2))

  #scrape data items
  systemname.auto <- website.auto %>%
    html_nodes("div+ table tr:nth-child(1) td:nth-child(2)") %>%
```

```

html_text()

PWSID.auto <- website.auto %>%
  html_nodes("td tr:nth-child(1) td:nth-child(5)") %>%
  html_text()

ownership.auto <- website.auto %>%
  html_nodes("div+ table tr:nth-child(2) td:nth-child(4)") %>%
  html_text()

MGD.auto <- website.auto %>%
  html_nodes("th~ td+ td") %>%
  html_text()

  #convert to a dataframe
df.auto <- data.frame("System Name" = systemname.auto,
                      "PWSID" = PWSID.auto,
                      "Ownership" = ownership.auto,
                      "MGD" = as.numeric(MGD.auto),
                      "Month" = c("Jan", "May", "Sep",
                                 "Feb", "Jun", "Oct",
                                 "Mar", "Jul", "Nov",
                                 "Apr", "Aug", "Dec"),
                      "Year" = rep(year2,12))

df.auto <- df.auto %>%
  mutate(Date = my(paste(Month, "-", Year))) %>%
  arrange(Date)

#return dataframe
return(df.auto)
}

```

7. Use the function above to extract and plot max daily withdrawals for Durham (PWSID='03-32-010') for each month in 2020. Then show the structure of the dataframe with either the `str()` or the `glimpse()` function.


```

#7
#use function to scrape Durham data
Durham.auto.df <- scrape.it('03-32-010',2020)

#plot dataframe
ggplot(Durham.auto.df, aes(
  x=Date,
  y=MGD,
))+
  geom_line()+
  labs(title = paste("2020 Max Daily Withdrawals in",
                    Durham.auto.df$System.Name))

```

2020 Max Daily Withdrawals in Durham



```
#show structure of dataframe  
str(Durham.auto.df)
```

```
## 'data.frame': 12 obs. of 7 variables:  
## $ System.Name: chr "Durham" "Durham" "Durham" "Durham" ...  
## $ PWSID      : chr "03-32-010" "03-32-010" "03-32-010" "03-32-010" ...  
## $ Ownership   : chr "Municipality" "Municipality" "Municipality" "Municipality" ...  
## $ MGD        : num 36 32 37.3 32.4 37 ...  
## $ Month       : chr "Jan" "Feb" "Mar" "Apr" ...  
## $ Year        : num 2020 2020 2020 2020 2020 2020 2020 2020 2020 2020 ...  
## $ Date        : Date, format: "2020-01-01" "2020-02-01" ...
```

8. Use the function above to extract data for Cary (PWSID = ‘03-32-010’) in 2020. Combine this data with the Durham data collected above and create a plot that compares Cary’s to Durham’s water withdrawals.

```
#8  
#use function to scrape Cary data  
Cary.auto.df <- scrape.it('03-92-020', 2020)  
  
#combine Durham and Cary dataframes  
Durham.Cary <- bind_rows(Durham.auto.df, Cary.auto.df)  
Durham.Cary
```

##	System.Name	PWSID	Ownership	MGD	Month	Year	Date
1	Durham	03-32-010	Municipality	36.0	Jan	2020	2020-01-01
2	Durham	03-32-010	Municipality	32.0	Feb	2020	2020-02-01
3	Durham	03-32-010	Municipality	37.3	Mar	2020	2020-03-01
4	Durham	03-32-010	Municipality	32.4	Apr	2020	2020-04-01
5	Durham	03-32-010	Municipality	36.0	May	2020	2020-05-01
6	Durham	03-32-010	Municipality	39.0	Jun	2020	2020-06-01
7	Durham	03-32-010	Municipality	43.5	Jul	2020	2020-07-01
8	Durham	03-32-010	Municipality	42.5	Aug	2020	2020-08-01
9	Durham	03-32-010	Municipality	42.5	Sep	2020	2020-09-01
10	Durham	03-32-010	Municipality	41.0	Oct	2020	2020-10-01
11	Cary	03-92-020	Private	29.0	Nov	2020	2020-11-01
12	Cary	03-92-020	Private	27.0	Dec	2020	2020-12-01

```

## 1 Durham 03-32-010 Municipality 36.01 Jan 2020 2020-01-01
## 2 Durham 03-32-010 Municipality 32.05 Feb 2020 2020-02-01
## 3 Durham 03-32-010 Municipality 37.29 Mar 2020 2020-03-01
## 4 Durham 03-32-010 Municipality 32.37 Apr 2020 2020-04-01
## 5 Durham 03-32-010 Municipality 36.98 May 2020 2020-05-01
## 6 Durham 03-32-010 Municipality 40.61 Jun 2020 2020-06-01
## 7 Durham 03-32-010 Municipality 43.63 Jul 2020 2020-07-01
## 8 Durham 03-32-010 Municipality 41.93 Aug 2020 2020-08-01
## 9 Durham 03-32-010 Municipality 41.69 Sep 2020 2020-09-01
## 10 Durham 03-32-010 Municipality 40.56 Oct 2020 2020-10-01
## 11 Durham 03-32-010 Municipality 33.32 Nov 2020 2020-11-01
## 12 Durham 03-32-010 Municipality 28.06 Dec 2020 2020-12-01
## 13 Cary 03-92-020 Municipality 21.95 Jan 2020 2020-01-01
## 14 Cary 03-92-020 Municipality 17.70 Feb 2020 2020-02-01
## 15 Cary 03-92-020 Municipality 20.56 Mar 2020 2020-03-01
## 16 Cary 03-92-020 Municipality 19.32 Apr 2020 2020-04-01
## 17 Cary 03-92-020 Municipality 23.22 May 2020 2020-05-01
## 18 Cary 03-92-020 Municipality 26.19 Jun 2020 2020-06-01
## 19 Cary 03-92-020 Municipality 28.01 Jul 2020 2020-07-01
## 20 Cary 03-92-020 Municipality 25.40 Aug 2020 2020-08-01
## 21 Cary 03-92-020 Municipality 25.49 Sep 2020 2020-09-01
## 22 Cary 03-92-020 Municipality 23.60 Oct 2020 2020-10-01
## 23 Cary 03-92-020 Municipality 21.90 Nov 2020 2020-11-01
## 24 Cary 03-92-020 Municipality 19.02 Dec 2020 2020-12-01


```

```

# create plot to compare Durham to Cary
ggplot(Durham.Cary, aes(
  x=Date,
  y=MGD,
  colour = System.Name
))+ 
  geom_line()+
  labs(title = paste("Max Daily Withdrawals"))

```

Max Daily Withdrawals



9. Use the code & function you created above to plot Cary's max daily withdrawal by months for the years 2018 thru 2023. Add a smoothed line to the plot (method = 'loess').

TIP: See Section 3.2 in the "10_Data_Scraping.Rmd" where we apply "map2()" to iteratively run a function over two inputs. Pipe the output of the map2() function to `bind_rows()` to combine the dataframes into a single one, and use that to construct your plot.


```
#9
#Cary max daily withdrawals years 2018-2023
PWSID2 = '03-92-020'
year2 = 2018:2023

the_dfs <- map2(PWSID2, year2, scrape.it) %>%
  bind_rows()

ggplot(the_dfs, aes(
  x=Date,
  y=MGD
)) +
  geom_line() +
  geom_smooth(method="loess", se=FALSE) +
  labs(
    title = paste("Max Daily Withdrawals in Cary")
  )

## `geom_smooth()` using formula = 'y ~ x'
```

Max Daily Withdrawals in Cary



Question: Just by looking at the plot (i.e. not running statistics), does Cary have a trend in water usage over time? > Answer: >Yes, there is an increase in the max daily withdrawals over time. There also >seems to be a seasonal pattern each year.